

Neural architecture search with target hardware control

Firsov Sergey

Supervisor: Oleg Bakhteev, PhD

Moscow Institute of Physics and Technology

2024

Problem

- ▶ **Neural Architecture Search:** Automate the design of neural network architectures to solve ML tasks
- ▶ **Goal:** Identify architectures that optimize task performance while balancing computational cost and resource efficiency.
- ▶ **Challenges:**
 - ▶ **Vast Search Space:** The number of possible architectures grows exponentially with model complexity.
 - ▶ **Trade-offs:** Balancing accuracy, model complexity, and hardware efficiency.
 - ▶ **Resource Constraints:** Models must perform well within predefined latency, memory, or energy limits.
 - ▶ **Multiple result:** Getting architectures as a function of complexity, thereby having multiple architectures in one answer.

Existing Methods

Exhaustive search for optimal architectures using:

- ▶ reinforcement learning
- ▶ evolutionary algorithms
- ▶ gradient-based methods


Differentiable Architecture Search (DARTS):

- ▶ formulates NAS as a continuous optimization problem, enabling gradient-based search.
- ▶ reduces computational cost by relaxing the discrete search space into a differentiable one.

Problems with Existing Methods:

- ▶ **accuracy vs. complexity:** often focus solely on accuracy, neglecting the importance of model complexity control.
- ▶ **hardware constraints ignored:** lack of mechanisms to account for real-world hardware constraints like latency or energy consumption.
- ▶ **limited flexibility:** insufficient control over the trade-off between accuracy and resource efficiency.

Architecture of solution



Step 1: Complexity-Aware Architecture Representation

- ▶ Architectural parameters α - functions from complexity parameter λ to architectural space \mathcal{A} :

$$\alpha : \lambda \rightarrow \mathcal{A}.$$

- ▶ Generates multiple architectures with varying complexities in a single optimization process.
- ▶ Ensures flexibility in architecture design across resource constraints.

Step 2: Simplex-Based Complexity Control

- ▶ Replace scalar complexity parameter (λ) with a simplex representation:

$$\mathbf{S} \in \Delta^{k-1}.$$

This makes complexity management more flexible for different types of operations.

- ▶ Δ^{k-1} : $(k - 1)$ -dimensional simplex, where k is the number of operations: now we have a different regularization coefficient for each operation.
- ▶ Sampling via Gumbel-Softmax enables differentiable selection of operations:

$$\mathbf{S} \sim \text{GumbelSoftmax}(\boldsymbol{\theta}).$$

Step 3: Hardware-Aware Latency Regularization

- ▶ **Latency-Aware Optimization:** Introduce a term in the loss function to account for operation latency on target hardware.

$$\text{Loss} = L + \kappa \cdot \text{Latency}(\alpha(\mathbf{S})),$$

where κ is a regularization coefficient.

- ▶ **Purpose:**
 - ▶ Ensure the resulting architecture meets real-world hardware constraints, such as inference time limits.
 - ▶ Optimize architectures not just for accuracy, but for deployment feasibility on specific devices.

Final Loss Function

- ▶ **Optimization problem:** minimize functional

$$\mathbb{E}_{\mathbf{S} \sim \text{GumbelSoftmax}} \left[L(\mathbf{w}^*(\mathbf{S}), \alpha(\mathbf{S})) + \lambda \cdot \text{Reg}(\alpha(\mathbf{S})) + \kappa \cdot \text{Latency}(\alpha(\mathbf{S})) \right].$$

- ▶ **Key Insights:**

- ▶ by using complexity dependence architecture parameters we generate multiple architectures in a single optimization process,
- ▶ by replacing the scalar complexity parameter with the simplex representation, we achieve finer control over architectural complexity, enabling flexible optimization across a range of hardware and resource constraints,
- ▶ by using latency lookup table, we can control hardware constraints.

Planned Experiments

- ▶ Validate the proposed method on:
 - ▶ Image classification benchmarks (e.g., fmnist, CIFAR-10).
 - ▶ Hardware-specific performance metrics.
- ▶ Compare against:
 - ▶ Baseline methods (DARTS, ProxylessNAS, FBNet).
 - ▶ Metrics: Accuracy, latency, and complexity.
- ▶ Evaluate trade-offs between accuracy and latency.

- ▶ Placeholder for experimental results.
- ▶ Future work: Analyze results and provide insights.

References I